

# amazon SQL

## Interview questions for Data Analysts

Part I



[linkedin.com/in/ileonjose](https://www.linkedin.com/in/ileonjose)



## 1. Average Review Ratings per Product per Month

- **Problem Statement:** Given a table of product reviews, calculate the average review rating for each product for every month. The data is in a reviews table, which includes *review\_id*, *user\_id*, *submit\_date*, *product\_id*, and *stars*. The output should list the month (as a numerical value), *product\_id*, and the average star rating rounded to two decimal places. Sort the result by month and then by *product\_id*.



review_id	user_id	submit_date	product_id	stars
6171	123	06/08/2022 0:00:00	50001	4
7802	265	06/10/2022 0:00:00	69852	4
5293	362	06/18/2022 0:00:00	50001	3
6352	192	07/26/2022 0:00:00	69852	3
4517	981	07/05/2022 0:00:00	69852	2



## How to Solve:

- Extract the month from the *submit\_date* using the *EXTRACT* function.
- Group the results by the extracted month and *product\_id*.
- Compute the average star rating for each group and round the result to two decimal places.
- Order the output by month and *product\_id*.



SQL

```
SELECT
    EXTRACT(MONTH FROM submit_date) AS mth,
    product_id,
    ROUND(AVG(stars), 2) AS avg_stars
FROM reviews
GROUP BY EXTRACT(MONTH FROM submit_date), product_id
ORDER BY mth, product_id;
```

[linkedin.com/in/ileonjose](https://www.linkedin.com/in/ileonjose)



## 2. Optimizing a Slow SQL Query

- **Problem Statement:** Amazon handles massive datasets, and optimizing SQL queries is crucial for performance. Discuss various methods to optimize a slow SQL query.



## How to Solve:

- Select Specific Fields: Use *SELECT field1, field2* instead of *SELECT \** to retrieve only necessary columns.
- Avoid SELECT DISTINCT: Use *DISTINCT* only when absolutely needed, as it can be expensive.
- Use INNER JOIN: Prefer *INNER JOIN* over using multiple *WHERE* clauses to join tables.



- Minimize Joins: Where possible, denormalize the data to reduce the need for complex joins.
- Add Indexes: Create indexes on columns that are frequently used in *WHERE* clauses and joins to speed up queries.
- Examine Execution Plans: Use the SQL query execution plan to identify bottlenecks and optimize accordingly.





### 3. SQL Constraints

- **Problem Statement:** Explain SQL constraints and provide examples of different types of constraints used to enforce data integrity in databases.



## How to Solve:

- NOT NULL: Ensures that a column cannot have NULL values.
- UNIQUE: Ensures all values in a column are unique.
- INDEX: Improves query performance by indexing frequently queried columns.
- PRIMARY KEY: Uniquely identifies each record in a table.
- FOREIGN KEY: Ensures referential integrity between tables.



SQL

```
CREATE TABLE employees (  
  employee_id INT PRIMARY KEY,  
  name VARCHAR(100) NOT NULL,  
  email VARCHAR(100) UNIQUE,  
  department_id INT,  
  FOREIGN KEY (department_id) REFERENCES departments(department_id)  
);
```

[linkedin.com/in/ileonjose](https://www.linkedin.com/in/ileonjose)



## 4. Highest-Grossing Items

- **Problem Statement:** Find the top two highest-grossing products in each category for the year 2022 from a table *product\_spend*. The table contains *category*, *product*, *user\_id*, *spend*, and *transaction\_date*. The output should include the *category*, *product*, and total spend.



category	product	user_id	spend	transaction_date
appliance	refrigerator	165	246.00	12/26/2021 12:00:00
appliance	refrigerator	123	299.99	03/02/2022 12:00:00
appliance	washing machine	123	219.80	03/02/2022 12:00:00
electronics	vacuum	178	152.00	04/05/2022 12:00:00
electronics	wireless headset	156	249.90	07/08/2022 12:00:00
electronics	vacuum	145	189.00	07/15/2022 12:00:00



## How to Solve:

- Step 1: Aggregate the total spend by category and product for 2022.
- Step 2: Use a Common Table Expression (CTE) to rank the products within each category based on total spend.
- Step 3: Filter the results to include only the top two products per category.



SQL

```
WITH product_category_spend AS (  
    SELECT  
        category,  
        product,  
        SUM(spend) AS total_spend  
    FROM product_spend  
    WHERE transaction_date >= '2022-01-01'  
        AND transaction_date <= '2022-12-31'  
    GROUP BY category, product  
,  
ranked_spend AS (  
    SELECT  
        category,  
        product,  
        total_spend,  
        RANK() OVER (PARTITION BY category ORDER BY total_spend DESC) AS  
ranking  
    FROM product_category_spend  
)  
SELECT  
    category,  
    product,  
    total_spend  
FROM ranked_spend  
WHERE ranking <= 2  
ORDER BY category, ranking;
```

[linkedin.com/in/ileonjose](https://www.linkedin.com/in/ileonjose)



## 5. Difference Between RANK() and DENSE\_RANK()

**Problem Statement:** Explain the difference between the RANK() and DENSE\_RANK() functions in SQL.





- `RANK()`: Assigns a unique rank to each row within a partition of a result set. If there are ties, the rank values will have gaps (e.g., if two items are ranked 2, the next rank will be 4).
- `DENSE_RANK()`: Similar to `RANK()`, but does not leave gaps between ranks. If two items are ranked 2, the next rank will be 3.



SQL

```
-- Using RANK()
SELECT
  product,
  sales,
  RANK() OVER (ORDER BY sales DESC) AS rank
FROM sales_data;

-- Using DENSE_RANK()
SELECT
  product,
  sales,
  DENSE_RANK() OVER (ORDER BY sales DESC) AS dense_rank
FROM sales_data;
```





Found this helpful? Repost!

**[linkedin.com/in/ileonjose](https://www.linkedin.com/in/ileonjose)**